

SALIH KILICLI – STAT 624 HOMEWORK 3

— SETUP —

`module load intel/2019`

I have written a C file that generates 1 billion exponentially distributed random numbers with $\alpha=2.0$, $\beta=3.0$ using different number generators and options. Lastly, I made the C file executable and then run it on my Terra account.

`chmod u+x hw3.c`

Problem 1) For this problem I have run the hw3.c file given below with Intel MKL using the commands

```
icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl  
  
time ./hw3.sh
```

Problem 2) I have compiled, link and run the code with default GNU C Compiler using the commands

```
gcc -o hw3 hw3.c -O2 -m64 -I${MKLROOT}/include /  
  
-L${MKLROOT}/lib/intel64 -Wl,-no-as-needed -lmkl_rt -lpthread -lm -ldl  
  
time ./hw3.sh
```

Problem 3) I have switched between options `VSL_BRNG_{250, MCG59, MT19937}` and run them using the same commands in problem 1.

- 1) `vslNewStream(&stream, VSL_BRNG_R250, 777);` // Generalized feedback shift register
- 2) `vslNewStream(&stream, VSL_BRNG_MCG59, 777);` // 59-bit multiplicative congruential
- 3) `vslNewStream(&stream, VSL_BRNG_MT19937, 777);` // Mersenne Twister pseudorandom

Problem 4) I have switched the fixed variables between `NVECTOR={10, 103, 106}` and `REPS={108, 106, 103}` to fix the # of random numbers 10⁹ [1 billion].

- | | | |
|----|---------------------------------------|-------------------------------|
| 1) | <code>#define NVECTOR 1000L</code> | <code>{10³}</code> |
| | <code>#define REPS 1000000L</code> | <code>{10⁶}</code> |
| 2) | <code>#define NVECTOR 10L</code> | <code>{10¹}</code> |
| | <code>#define REPS 100000000L</code> | <code>{10⁸}</code> |
| 3) | <code>#define NVECTOR 1000000L</code> | <code>{10⁶}</code> |
| | <code>#define REPS 1000L</code> | <code>{10³}</code> |

The outputs for the problems and comparisons between different options is given in Final Remarks.

SALIH KILICLI – STAT 624 HOMEWORK 3

— HW3.C FILE —

```
#include <stdio.h>

#include "mkl_vsl.h"

#include "mkl.h"

#define NVECTOR 1000L           //10L           //1000000L
#define REPS 1000000L         //100000000L       //1000L

// icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl           # to compile in Intel
// gcc -o hw3 hw3.c -O2 -xHost -m64 -I${MKLROOT}/include \           # to compile in GNU GCC
// -L${MKLROOT}/lib/intel64 -W1, -no-as-needed -lmkl_rt -lpthread -lm -ldl
// time ./hw3                                                         # to run and time

int main() // hw3.c
{
    /* Initializing r and average s */

    double r[NVECTOR], s=0.0;

    VSLStreamStatePtr stream;

    int i, j;

    //vslNewStream( &stream, VSL_BRNG_R250, 777 );           // A generalized feedback shift register generator
    //vslNewStream( &stream, VSL_BRNG_MCG59, 777 );           // 59-bit multiplicative congruential generator
    vslNewStream( &stream, VSL_BRNG_MT19937, 777 );           // Mersenne Twister pseudorandom num generator

    /* Generating random numbers */

    for ( i=0; i<REPS; i++ ) {

        vdRngExponential(VSL_RNG_METHOD_EXPONENTIAL_ICDF, stream, NVECTOR, r, 2.0, 3.0 );

        for ( j=0; j<NVECTOR; j++ ) s += r[j];

    }
```

SALIH KILICLI – STAT 624 HOMEWORK 3

```
/* Deleting the stream */
```

```
vslDeleteStream( &stream );
```

```
printf( "Sample mean = %f (n=%ld)\n", s/(REPS*NVECTOR), REPS*NVECTOR );
```

```
return 0;
```

```
}
```

— FINAL REMARKS —

Running hw3.c executable file in Terra and collecting time information, I have created tables below in order to compare random generators and different options. I have also attached the outputs.

Problem 1: Output 1 – NVECTOR = 10^3 – REPS = 10^6 – Intel MKL with Exponential Distribution

```
[math3mantic@terra2 ~]$ ls
hw1 hw2 hw3 perl5
[math3mantic@terra2 ~]$ cd ~/hw3
[math3mantic@terra2 hw3]$ ls
hw3 hw3.c
[math3mantic@terra2 hw3]$ module load intel/2019
[math3mantic@terra2 hw3]$ icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl
[math3mantic@terra2 hw3]$ time ./hw3
Sample mean = 4.999862 (n=1000000000)

real    0m2.704s
user    0m2.478s
sys     0m0.031s
```

Problem 2: Output 2 – NVECTOR = 10^3 – REPS = 10^6 – GNU GCC

```
[math3mantic@terra1 hw3]$ gcc -o hw3 hw3.c -O2 -m64 -I${MKLROOT}/include -L${MKLROOT}/lib/intel64 -Wl,--no-as-needed -lmkl_rt -lpthread -lm -ldl
[math3mantic@terra1 hw3]$ time ./hw3
Sample mean = 4.999862 (n=1000000000)

real    0m4.015s
user    0m3.320s
sys     0m0.026s
```

Random Number Generator	Intel MKL	Linked GNU GCC
Mersenne Twister	real 0m2.701s user 0m2.684s sys 0m0.016s	real 0m4.015s user 0m3.320s sys 0m0.026s

SALIH KILICLI – STAT 624 HOMEWORK 3

It is clear from the outputs and the table given above that Intel MKL is almost 2x faster than default GNU C compiler in terms of speed of generating random numbers.

Problem 3: Output 3 – NVECTOR = 10^3 – REPS = 10^6 – Intel MKL with different generators

Outputs for 1) Generalized feedback shift operator, 2) 59-bit multiplicative congruential generator and the 3) Mersenne Twister pseudorandom number generator for Intel MKL given below.

```
[math3mantic@terra2 hw3]$ icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl
[math3mantic@terra2 hw3]$ time ./hw3
Sample mean = 4.999996 (n=1000000000)

real    0m3.322s
user    0m3.303s
sys      0m0.018s
[math3mantic@terra2 hw3]$ icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl
[math3mantic@terra2 hw3]$ time ./hw3
Sample mean = 4.999999 (n=1000000000)

real    0m2.390s
user    0m2.372s
sys      0m0.017s
[math3mantic@terra2 hw3]$ icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl
[math3mantic@terra2 hw3]$ time ./hw3
Sample mean = 4.999862 (n=1000000000)

real    0m2.724s
user    0m2.702s
sys      0m0.020s
```

Table below compares the speed of different random Intel MKL number generators:

Random Number Generator	Intel	
59-bit multiplicative congruential generator	real	0m2.390s
VSL_BRNG_MCG59	user	0m2.372s
Sample Mean = 4.999999	sys	0m0.017s
The Mersenne Twister pseudorandom generator	real	0m2.724s
VSL_BRNG_MT19937	user	0m2.702s
Sample Mean = 4.999862	sys	0m0.020s
Generalized feedback shift register generator	real	0m3.322s
VSL_BRNG_R250	user	0m3.303s
Sample Mean = 4.999996	sys	0m0.018s

SALIH KILICLI – STAT 624 HOMEWORK 3

It is evident from the table that for Intel RNG, fastest method is MCH59 which also gives the best estimate for sample mean.

Problem 4: Output 4 – Intel MKL number generator using Mersenne Twister is tested for different NVECTORS. Outputs for NVECTOR = { 10^3 , 10, 10^6 } and REPS = { 10^6 , 10^8 , 10^3 }, respectively:

```
[math3mantic@terra2 hw3]$ module load intel/2019
[[math3mantic@terra2 hw3]$ icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl
[[math3mantic@terra2 hw3]$ time ./hw3
Sample mean = 4.999862 (n=1000000000)

real    0m2.704s
user    0m2.478s
sys     0m0.031s
[[math3mantic@terra2 hw3]$ icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl
[[math3mantic@terra2 hw3]$ time ./hw3
Sample mean = 4.999862 (n=1000000000)

real    0m21.274s
user    0m21.078s
sys     0m0.030s
[[math3mantic@terra2 hw3]$ icc -o hw3 hw3.c -O2 -xHost -lmkl_rt -lpthread -lm -ldl
[[math3mantic@terra2 hw3]$ time ./hw3
Sample mean = 4.999862 (n=1000000000)

real    0m2.970s
user    0m2.946s
sys     0m0.022s
```

Intel MKL RNG	N=10		N=1000		N=1000,000	
Mersenne Twister	real	0m21.274s	real	0m2.704s	real	0m2.970s
Sample Mean=4.999862	user	0m21.078s	user	0m2.478s	user	0m2.946s
	sys	0m0.030s	sys	0m0.031s	sys	0m0.022s

When the number of random numbers generated at a time is 10, the time to run increases dramatically. However, when the number of random numbers generated at a time is 1000 it takes less time to run than 1000,000 at a time. It is apparently because of the CPU architecture and neither generating very small nor very big numbers at a time yield the best result in terms of speed.